

Top 10 LLM – Details

Model = LLM name/provider

Release date = public/API release

Cutoff date = model's stated knowledge/training cutoff

Cost/token = API price per 1M input/output tokens

Context window = max tokens the model can read in one request

Best use cases = where the model fits best

#	Model	Release date	Cutoff date	Cost / 1M tokens	Context window	Best use cases	One-line statement
1	OpenAI GPT-5.5	Apr 23, 2026	Dec 1, 2025	\$5 input / \$30 output	1,050,000	Complex coding, research, data analysis, document-heavy work	Best fit when quality and long-context professional reasoning matter most. (OpenAI)
2	Anthropic Claude Opus 4.7	Apr 16, 2026	Jan 2026	\$5 input / \$25 output	1,000,000	Senior-level coding, agents, enterprise workflows	Strong premium choice for hard agentic and software-engineering tasks. (anthropic.com)
3	Google Gemini 3.1 Pro Preview	Feb 19, 2026	Jan 2025	\$2/\$12 under 200k; \$4/\$18 over 200k	1,048,576 input / 65,536 output	Multimodal reasoning, PDFs, code repos, finance/spreadsheets agents	Great for huge multimodal inputs and Google ecosystem workflows. (Google Cloud Documentation)
4	Anthropic Claude Sonnet 4.6	2026 generation	Aug 2025; training cutoff Jan 2026	\$3 input / \$15 output	1,000,000	Balanced coding, agents, production apps	Best default Claude-style balance of quality, speed, and cost. (platform.claude.com)

#	Model	Release date	Cutoff date	Cost / 1M tokens	Context window	Best use cases	One-line statement
5	OpenAI GPT-5.1	Nov 13, 2025	Sep 30, 2024	\$1.25 input / \$10 output	400,000	General apps, coding, agentic workflows, structured outputs	Practical all-rounder when GPT-5.5 quality is not required. (OpenAI Developers)
6	Qwen3-Max	Sep 23, 2025	Jun 30, 2025	\$0.78 input / \$3.90 output	262,144	Multilingual apps, coding, RAG, tool calling	Excellent value model for Chinese-English, RAG, and scalable enterprise tasks. (OpenRouter)
7	Mistral Medium 3.1	Aug 12/13, 2025	Jun 30, 2025 via OpenRouter metadata	\$0.40 input / \$2 output	128k–131k	Enterprise chat, STEM, coding, multimodal apps	Cost-efficient frontier-style model for European/enterprise deployments. (Mistral AI Documentation)
8	xAI Grok 4	Jul 9, 2025	Official xAI page says Nov 2024; third-party API page says Jul 2025	\$3 input / \$15 output	256k	Real-time/web-heavy apps, reasoning, tool use	Best when X/web-connected reasoning is part of the product experience. (xAI Docs)
9	Moonshot Kimi K2.6	Apr 21, 2026	Jan 2025	\$0.95 input / \$4 output	262,144	Agentic tasks, research, software development, long conversations	Strong open-weight-style option for agentic research and coding workflows. (TypingMind)
10	DeepSeek V4 Pro / Flash	Apr 2026 preview	Not officially published in docs found	Flash: \$0.14 input / \$0.28 output;	1,000,000	Low-cost reasoning, coding, large-context Chinese/English workloads	Very low-cost large-context option, but cutoff transparency is weaker. (DeepSeek API Docs)

Below are **practical examples of what each model is good at**.

#	Model	Example use case
1	OpenAI GPT-5.1	"Build a Python API, write tests, explain bugs, and generate clean JSON output for an automation workflow." OpenAI lists GPT-5.1 as a flagship model for coding and agentic tasks with configurable reasoning. (OpenAI Developers)
2	Claude Opus 4.5	"Analyze a large codebase, identify architectural issues, create a refactor plan, and update files step by step." Anthropic describes Opus 4.5 as strong for coding, agents, computer use, research, slides, and spreadsheets. (anthropic.com)
3	Gemini 3 Pro / 3.1 Pro	"Upload a 300-page policy PDF plus charts and ask it to summarize risks, compare clauses, and produce an executive brief." Gemini 3 supports a 1M-token input context and up to 64k output tokens. (Google AI for Developers)
4	Claude Sonnet 4.5	"Create a customer-support agent that reads tickets, calls tools, drafts replies, and escalates uncertain cases." Claude Sonnet 4.5 pricing is listed at \$3 input / \$15 output per million tokens. (anthropic.com)
5	Qwen3-Max	"Build a bilingual chatbot for English and Chinese users that searches internal documents and calls business APIs." Qwen3-Max is listed with 262k context, tool-calling support, and a Jun 30, 2025 knowledge cutoff by OpenRouter. (OpenRouter)
6	Mistral Medium 3.1	"Create a private enterprise assistant that answers employee questions from policies, emails, and technical docs at lower cost." Mistral's model card lists Medium 3.1 as a frontier-class multimodal model with 128k context and \$0.40/\$2 per million token pricing. (Mistral AI Documentation)
7	xAI Grok 4	"Create an assistant that explains trending topics, compares public opinions, and drafts responses for fast-moving discussions." xAI's docs list Grok 4 knowledge cutoff as November 2024, while third-party API listings show Grok 4 0709 with 256k context and \$3/\$15 pricing. (xAI Docs)
8	Kimi K2.6	"Give it a full project repo and ask it to find bugs, propose patches, and generate migration documentation." Kimi K2.6 is listed with Apr 21, 2026 release, 262,144-token context, and \$0.95/\$4 per million token pricing. (TypingMind)
9	DeepSeek V4	"Run large-scale code review or document Q&A where cost per request matters more than premium polish." Recent reports describe DeepSeek V4 as having a 1M-token context window and low-cost Pro/Flash variants, though official cutoff details are less transparent. (Tom's Hardware)
10	GPT-4.1	"Paste a very large repo or documentation set and ask for bug fixes, code generation, and implementation guidance." GPT-4.1 was released in April 2025 with up to 1M-token context and improved coding/long-context performance. (Reuters)